


Same architecture, seed, optimizer and step budget; different training positions

Reconstruction loss on training batches



Same full-trace dev sample

